
PRETTY DARN GOOD CONTROL: WHEN ARE APPROXIMATE SOLUTIONS BETTER THAN APPROXIMATE MODELS

A PREPRINT

Felipe Montealegre-Mora Marcus Laperolerie Melissa Chapman Abigail G. Keller

Carl Boettiger

cboettig@berkele.edu

January 17, 2023

Abstract

Write

Keywords Optimal Control · Reinforcement Learning · Uncertainty · Decision Theory

1 Introduction

Despite the complexity inherent in natural ecosystems, analytically tractable models have long played an essential role in their understanding and management. Ecological management may be understood as an optimal control problem - a manager must determine a sequence of actions (a “policy”) to optimize some desired objective(s) and constraints (e.g. conservation goals, natural resource extraction, economic and social constraints.)

2 Figures brainstorm

Figure 0: contrast 1D vs 3D strategies (see commented out diagram).

11 Figure 1: 1-D and 3-D model conceptual figure. something about the objective / decision Figure 2: Stability
 12 / multistability in the 1D and 3D models. state space + time views Figure 3: The 1-D optimal management
 13 solution. ‘constant escapement’ intuition etc
 14 RL figures: - schematic of RL (a-la previous marcus paper Fig 1) - Neural network optimization figure
 15 Results figures: - timeseries example of management under the RL policy. probably compare to managing
 16 under 1-D solution / rule-of-thumb methods - state-space view of management doughnut - visualization /
 17 encapsulation of the policy heatmaps, slices, policy vs position along ellipse - reward plot over time (comparing
 18 methods)

19 3 Mathematical models of fisheries

In this section we introduce some models describing the population dynamics of fisheries. In general, the class of models that appear in this context are *first order finite difference equations*. For n species, these models have the general form

$$N_{t+1} - N_t = f(N_t) - h_t := \Delta N_t, \quad (1)$$

20 where $N_t = (X_t, Y_t, \dots) \in \mathbb{R}_+^n$ is a vector of populations, $f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ is an arbitrary function, and h_t is
 21 the harvest size at timestep t .

22 3.1 A single species model

23 Most commonly, fisheries use 1-dimensional models, $n = 1$. One such model appearing in a variety of
 24 ecological contexts is *logistic growth*, for which

$$f(X_t) = rX_t(1 - X_t/K) =: L(X_t; r, K).$$

More relevant to our current work, is the following single-species model studied in [1],

$$\Delta X_t = L(X_t; r, K) - F(X_t, H; \beta, c) - h_t, \quad (2)$$

where,

$$F(X, H; \beta, c) := \frac{\beta H X^2}{c^2 + X^2}.$$

25 The model has five parameters: the growth rate r and carrying capacity K for X , a constant population H
 26 of a species which preys on X , the maximal rate of predation β , and the predation half-maximum biomass c .
 27 Eq. (2) is an interesting study case of a *tipping point* (see Fig. ??). Holding the value of β fixed, for
 28 intermediate values of H there exist two stable fixed points for the state V_1 of the system, these two attractors

separated by an unstable fixed point. At a certain threshold value of H , however, the top stable fixed point collides with the unstable fixed point and both are annihilated. For this value of H , and for higher values, only the lower fixed point remains.

This structure implies two things. First, that a drift in H could lead to catastrophic consequences, with the population V_1 plummeting to the lower fixed stable point. Second, that if the evolution of V_1 is *stochastic*, then, even at values of H below the threshold point, the system runs a sizeable danger of tipping over towards the lower stable point.

3.2 A three species model

Consider the following three species *generalization* of eq. (2):

$$\begin{aligned}\Delta X_t &= L(X_t; r_X, K_X) - F(X_t, Z_t; \beta, c) - c_{XY}X_tY_t - h_t, \\ \Delta Y_t &= L(Y_t; r_Y, K_Y) - DF(Y_t, Z_t; \beta, c) - c_{XY}X_tY_t, \\ \Delta Z_t &= (f(X_t + DY_t) - d_Z)Z_t.\end{aligned}\tag{3}$$

The model is conceptualized in Fig. ?? . It contains three populations, the state at time t is (X_t, Y_t, Z_t) . Species Z preys on both X and Y , while the latter two compete for resources. There are ten parameters in this model: The growth rate and carrying capacity, r_X , K_X , r_Y and K_Y , of X and Y . A parameter c_{XY} mediating a Lotka-Volterra competition between X and Y . A maximum hunting intensity β and half-maximum c specifying how Z forages on X and Y . A parameter D regulating a relative preference of Z to prey on Y . Finally, a death rate d_Z and a parameter f related to the birth rate of Z .

This model generalizes eq. (2) in the following sense: If $f = d_Z = c_{XY} = 0$, and we let $Z_{t=0} = H$, then eq. (3) becomes

$$\begin{aligned}\Delta X_t &= L(X_t; r_X, K_X) - F(X_t, H; \beta, c) - h_t, \\ \Delta Y_t &= L(Y_t; r_Y, K_Y) - DF(Y_t, H; \beta, c), \\ Z_t &= H.\end{aligned}$$

The equations for X_t and Y_t thus become decoupled, with X_t evolving according to (2).

As can be expected, the dynamics of eq. (3) are significantly more complex than those of eq. (2). There are a variety of fixed points for the equations—around 7 or 8 for wide ranges of parameter values. Fig. ?? provides a rough idea of the fixed point structure by plotting the bifurcation diagram obtained from varying D and fixing the other parameters.

Throughout most of this paper we will fix the parameters of the model according to Table ?? . At these values, the fixed points of the model are:

$$TBD$$

The classification of fixed points for multi-dimensional systems is generally more complicated than in one dimension. A fixed point might be stable against perturbations in one direction, but unstable against perturbations along another. While a detailed classification of these fixed points is beyond the scope of this work, we point out that the point (*TBD*) is stable against perturbations in all directions. We will thus use it as a starting point for our simulations.

3.3 Stochastic dynamics

Probably just merge on previous sections. We consider a stochastic contribution to the system's dynamics. This way, eq. (2) becomes

$$\Delta X_t = L(X_t; r, K) - F(X_t, H; \beta, c) - h_t + \eta_t, \quad (4)$$

where $\{\eta_t\}_t$ are independent Gaussian variables with zero mean. The variance σ^2 of η_t is an additional parameter of the model.

Similarly, eq. (3) becomes

$$\begin{aligned} \Delta X_t &= L(X_t; r_X, K_X) - F(X_t, Z_t; \beta, c) - c_{XY} X_t Y_t - h_t + \eta_{X,t}, \\ \Delta Y_t &= L(Y_t; r_Y, K_Y) - DF(Y_t, Z_t; \beta, c) - c_{XY} X_t Y_t + \eta_{Y,t}, \\ \Delta Z_t &= (f(X_t + DY_t) - d_Z) Z_t + \eta_{Z,t}. \end{aligned}$$

4 Fishery management approaches

Here we review classical strategies for sustainable fishery management and we provide a birds-eye view of the alternative approach we propose. A summary of the contrast between the two strategies is given in Fig. ??.

4.1 Classical approaches to sustainable fisheries

There are several strategies that have been used to manage fisheries: *escapement*, *maximum sustainable yield (MSY)*, *total allowable catch (TAC)*, among others. We collectively refer to these as *classical*, and will compare their performance to RL-based management strategies. As shown in Fig. ??, classical strategies have the common aspect of reducing the complex dynamics of the fishery ecosystem to a single equation governing the harvested population (say, F). A common example is using a logistic growth equation,

$$F_{t+1} - F_t = rF_t(1 - F_t/K) - h_t =: L(F_t) - h_t,$$

where the interaction between F and its environment is summarized to two parameters, the growth rate r , and the carrying capacity K . In the equation above, h_t is the *harvest* at timestep t . The goal is to choose the harvest policy $h : F_t \mapsto h_t$ such that long-term profits are maximized.

An advantage of one dimensional approaches is that the optimal policy is often known exactly, and, moreover, is intuitive. For example, in the logistic equation pointed out above, the maximal sustainable yield of the system is attained at $F = F_{MSY} := K/2$. The optimizer is an *escapement* policy:

$$h_t = \begin{cases} F_t - K/2, & \text{if, } F_t > K/2 \\ 0, & \text{else.} \end{cases}$$

This corresponds to keeping the system at its optimal growth rate as much as possible.

Escapement policies, or more generally *bang bang* policies, tend to be the optimal solution for these types of control problems. A drawback of these solutions, in the fishery context, is the presence of several timesteps with zero harvest which can arise. To mend this, certain suboptimal solutions have been constructed for fishery management.

One ubiquitous solution is simply called maximum sustainable yield (MSY). It consists on letting $h(F) = rF/2$, so that

$$h(F_{MSY}) = L(F_{MSY}) = rK/4.$$

That is, at the MSY biomass, the logistic growth of F is cancelled exactly by the harvest.

The MSY rule fixes the drawback in the escapement policy by having $h(F) > 0$ for all $F > 0$. It, however, has its own drawbacks. It is particularly sensitive to misestimates of the parameter r , as we will discuss in Sec. ???. Due to this, similar but more conservative policies have been used.

Total allowable catch (TAC) is one such policy. It consists on reducing the inclination of the line defined by $h(F)$ using a prefactor α in $h(F) = \alpha \times rF/2$. Plausible examples are $\alpha = 0.8$ or 0.9 . Another common alternative is to have a prefactor $\alpha = \alpha(F)$ which decreases from one to zero as F decreases below a threshold:

$$h(F) = \alpha(F) \times rF/2,$$

where,

$$\alpha(F) = \begin{cases} 1, & \text{if } F > F_{\text{thresh.}}, \\ F/F_{\text{thresh.}}, & \text{else.} \end{cases}$$

We call this policy *variable total allowable catch (VTAC)*.

4.2 Dangers: Noisy parameters, noisy observations and model inaccuracy

5 Reinforcement learning

Reinforcement learning (RL) is, in a nutshell, a way of approaching *control problems* through machine learning. An RL algorithm can be conceptually separated into two parts: an *agent*, and an *environment* which the agent can interact with. That is, the agent may act on the environment and thus change its state, while the environment gives a *reward* to the agent in return (see Fig. ??). The rewards encode the agent’s goal. The main part of an RL algorithm is then to progressively improve the agent’s *policy*, in order to maximize the cumulative reward received. This is done by aggregating experience and learning from it.

For our use case, the environment will be a computational model of the population dynamics of a fishery, with the environment state being a vector of all the fish populations, $S = (V_1, V_2, H)$. At each time step, the agent harvests one of the populations, V_1 . This changes the state as

$$(V_1, V_2, H) \mapsto ((1 - q)V_1, V_2, H),$$

where q is a fishing *quota* set by the agent. This secures a reward of qV_1 . Afterwards, the environment undergoes a timestep under its natural dynamics given by (3).

5.1 Mathematical framework for RL

Mathematically, RL may be formulated using a discrete time *partially observable Markov decision process* (POMDP). This formalization is rather flexible and allows one, e.g., to account for situations where the agent may not fully observe the environment state, or where the only observations available to the agent are certain functions of the underlying state. For the sake of clarity, we will present here only the class of POMDPs which are relevant to our work: *fully observable MDPs with a trivial emission function* (FMDPs for short). An FMDP may be defined by the following data:

- \mathcal{S} : *state space*, the set of states of the environment,
- \mathcal{A} : *action space*, the set of actions which the agent may choose from,
- $T(s_{t+1}|s_t, a_t)$: *transition operator*, a conditional distribution which describes the dynamics of the system,
- $r(s_t, a_t)$: *reward function*, the reward obtained after performing action $a_t \in \mathcal{A}$ in state s_t ,
- $d(s_0)$: *initial state distribution*, the initial state of the environment is sampled from this distribution,
- $\gamma \in [0, 1]$: *discount factor*.

At a time t , the FMDP agent observes the full state s_t of the environment and chooses an action based on this observation according to a *policy function* $\pi(a_t|s_t)$. In return, it receives a discounted reward $\gamma^t r(a_t, s_t)$.

The discount factor helps regularize the agent, helping the optimization algorithm find solutions which pay off within a timescale of $t \sim \log(\gamma^{-1})^{-1}$.

With any fixed policy function, the agent will traverse a path $\tau = (s_0, a_0, s_1, a_1 \dots, s_{t_{\text{fin.}}})$ sampled randomly from the distribution

$$p_\pi(\tau) = d(s_0) \prod_{t=0}^{t_{\text{fin.}}-1} \pi(a_t|s_t)T(s_{t+1}|s_t, a_t).$$

Reinforcement learning seeks to optimize π such that the expected rewards are maximal,

$$\pi^* = \operatorname{argmax} \mathbb{E}_{\tau \sim p_\pi}[R(\tau)],$$

where,

$$R(\tau) = \sum_{t=0}^{t_{\text{fin.}}-1} \gamma^t r(a_t, s_t),$$

is the cumulative reward of path τ . The function $J(\pi) := \mathbb{E}_{\tau \sim p_\pi}[R(\tau)]$ is called the *value function*.

5.2 Deep Reinforcement Learning

The policy function π often lives in a high- or even infinite-dimensional space. This makes it unfeasible to directly optimize π . In practice, an alternative approach is used: π is optimized over a much lower-dimensional parametrized family of functions. Deep reinforcement learning uses this strategy, focusing on function families parametrized by neural networks. (See Fig. ??.)

Since our state and observation spaces are continuous, we will focus on deep reinforcement learning throughout this paper. Specifically, we parametrize π using a neural network with two hidden layers of 64 neurons.

We use the *proximal policy optimization (PPO)* algorithm to optimize π . Within the RL literature there is a wealth of algorithms from which to choose from, each with its pros and cons. Here we have focused on a single one of these for the sake of particularity—to draw a clear comparison between the RL-based and the classical fishery management approaches. In practice, further improvements can be expected by a careful selection of the optimization algorithm.

5.3 Model-free reinforcement learning

Within control theory, the usual setup is one where we use as much information from the model as possible in order to derive an optimal solution. **(Wax poetic a bit about Bellmann eq. approaches here?)**

The classical sustainable fishery management approaches summarized in Sec. 4.1 are model-based controls. As we saw in that section, these controls may run into trouble in the case where there are inaccuracies in the model parameter estimates.

More generally, there are many situations in which the exact model of the system is not known or not tractable. This is a standard situation in ecology: mathematical models capture the most prominent aspects of the ecosystem’s dynamics, while ignoring or summarizing most of its complexity. In this case, it is clear, model-based controls run a grave danger of mismanaging the system.

Reinforcement learning is, on the other hand, typically a model-free approach to control theory.¹ While the model is certainly used to generate training data, it is not directly used by the optimization algorithm. This provides more flexibility to use model-free RL in instances where the model of the system is not accurately known. In fact, it has been shown to generally be preferable to use model-free RL in such instances.

Mention CL here? Where solutions to other related problems in the curriculum are used to build solutions to the target problem.

This context provides a motivation for this paper. Indeed, models for ecosystem dynamics are only ever approximate and incomplete descriptions of reality. This way, it is plausible that model-free RL controls outperform currently used model-based controls in ecological management problems.

6 Results

We trained a DRL agent to obtain a candidate policy

$$\pi : (X_t, Y_t, Z_t) \mapsto h_t$$

for eq. (3). The agent could fully and precisely observe the state of the system at each time step and decide on its harvest based on this information. The agent was trained using the PPO algorithm, using a neural network with two hidden layers of 64 neurons to express π . Details of the training process may be found in App. ???. Our code is available at TBD.

We compare the performance of this agent to the performance of MSY and of an 80% TAC.

Our results suggest that DRL can find solutions that are more sustainable *and* more profitable than both MSY and TAC. Indeed, DRL finds solutions in which, for certain periods of time, no harvest is performed on the system. These periods effectively allow the system to, in a sense, “reset” itself. They lead to a distinctive oscillating behavior of the system state over time, with periods of relative depletion followed by periods of growth. Figure ?? exemplifies this point by comparing the harvest effort and state times series obtained from these policies.

¹There are, however, approaches to perform model-based reinforcement learning. While we will not focus on these in this paper they are discussed in [?, ?].

It is well known that following MSY policies can lead to over-fishing. Our results reproduce this behavior and, furthermore, warn that using a TAC policy instead might not be enough to gain sustainability. Figure TBD-A shows the mean episode length for

6.1 Acknowledgements

The title of this piece references a mathematical biology workshop at NIMBioS organized by Paul Armsworth, Alan Hastings, Megan Donahue, and Carl Tows in 2011 that first posed the question addressed here.

7 References

References

- [1] Robert M May. Thresholds and breakpoints in ecosystems with a multiplicity of stable states. *Nature*, 269(5628):471–477, 1977.